

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (confusion_matrix, classification_report, accuracy_score, r2_score)
```

```
In [4]: df = pd.read_csv('Social_Network_Ads.csv')
```

```
In [5]: df.head()
```

```
Out[5]:
```

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|--|---------|--------|-----|-----------------|-----------|
|--|---------|--------|-----|-----------------|-----------|

| | | | | | |
|---|----------|--------|----|-------|---|
| 0 | 15624510 | Male | 19 | 19000 | 0 |
| 1 | 15810944 | Male | 35 | 20000 | 0 |
| 2 | 15668575 | Female | 26 | 43000 | 0 |
| 3 | 15603246 | Female | 27 | 57000 | 0 |
| 4 | 15804002 | Male | 19 | 76000 | 0 |

```
In [6]: df.shape
```

```
Out[6]: (400, 5)
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
In [11]: print("\n=== Statistical Information ===")
df.describe()
```

```
=== Statistical Information ===
```

```
Out[11]:
```

| | User ID | Age | EstimatedSalary | Purchased |
|-------|--------------|------------|-----------------|------------|
| count | 4.000000e+02 | 400.000000 | 400.000000 | 400.000000 |
| mean | 1.569154e+07 | 37.655000 | 69742.500000 | 0.357500 |
| std | 7.165832e+04 | 10.482877 | 34096.960282 | 0.479864 |
| min | 1.556669e+07 | 18.000000 | 15000.000000 | 0.000000 |
| 25% | 1.562676e+07 | 29.750000 | 43000.000000 | 0.000000 |
| 50% | 1.569434e+07 | 37.000000 | 70000.000000 | 0.000000 |
| 75% | 1.575036e+07 | 46.000000 | 88000.000000 | 1.000000 |
| max | 1.581524e+07 | 60.000000 | 150000.000000 | 1.000000 |


```
In [10]: df.isnull().sum()
```

```
Out[10]: User ID      0
Gender      0
Age         0
EstimatedSalary  0
Purchased   0
dtype: int64
```

```
In [12]: # Drop User ID as it's not relevant
df = df.drop('User ID', axis=1)
```

```
In [13]: # Convert Gender to numerical values (0 for Female, 1 for Male)
df['Gender'] = df['Gender'].map({'Female': 0, 'Male': 1})
```

```
In [14]: # Display correlation matrix
plt.figure(figsize=(8,6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
In [15]: # Prepare features and target
X = df[['Age', 'EstimatedSalary']]
y = df['Purchased']

# Split data into training and testing sets (75% train, 25% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
In [16]: # Build and train Logistic regression model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

Out[16]:

▼ LogisticRegression ⓘ ?

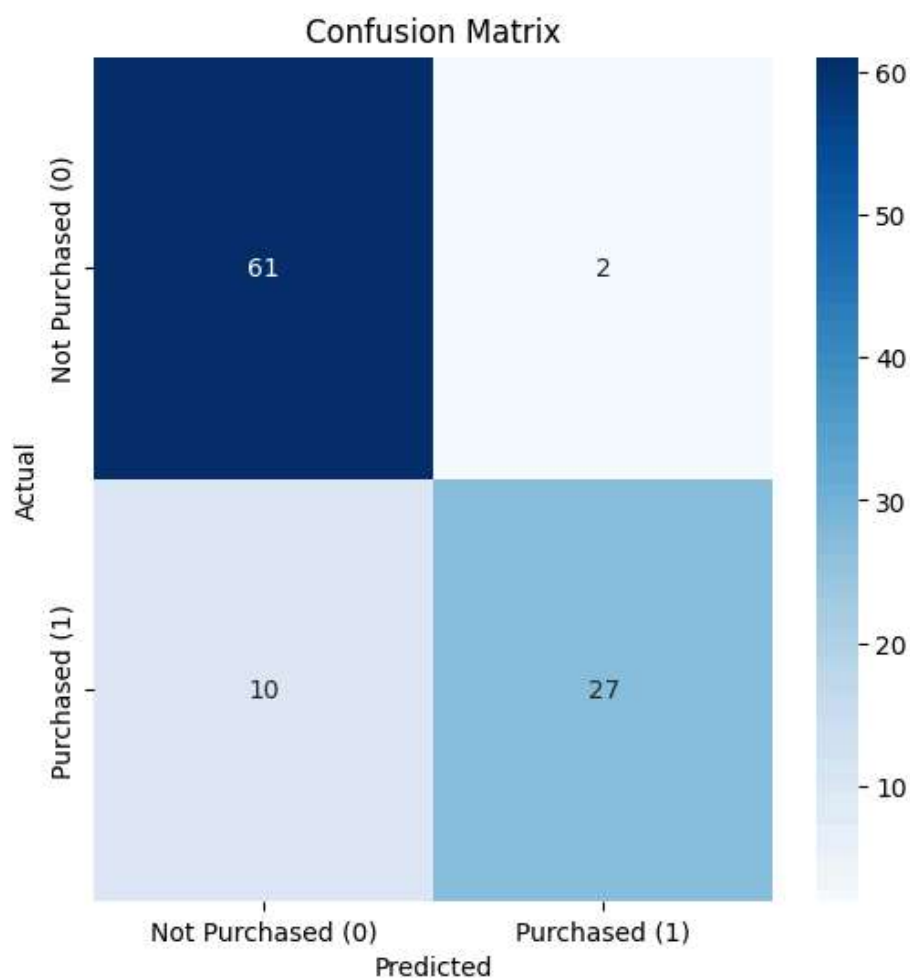
LogisticRegression()

```
In [17]: # Make predictions
y_pred = logreg.predict(X_test)
y_prob = logreg.predict_proba(X_test)[:, 1]
```

```
In [18]: cm = confusion_matrix(y_test, y_pred)
print("\n=== Confusion Matrix ===")
print(cm)
```

```
=== Confusion Matrix ===
[[61  2]
 [10 27]]
```

```
In [19]: # Visualize confusion matrix
plt.figure(figsize=(6,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
             xticklabels=['Not Purchased (0)', 'Purchased (1)'],
             yticklabels=['Not Purchased (0)', 'Purchased (1)'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
In [22]: # Calculate metrics
tn, fp, fn, tp = cm.ravel()
accuracy = (tp + tn) / (tp + tn + fp + fn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
```

```
recall = tp / (tp + fn)
r2 = r2_score(y_test, y_prob)
```

```
In [23]: # Display classification report and metrics
print("\n=== Classification Report ===")
print(classification_report(y_test, y_pred))
print("\n=== Performance Metrics ===")
print(f"True Positives (TP): {tp}")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")
print(f"Accuracy: {accuracy:.4f}")
print(f"Error Rate: {error_rate:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"R² Score (using probabilities): {r2:.4f}")
```

```
=== Classification Report ===
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.86 | 0.97 | 0.91 | 63 |
| 1 | 0.93 | 0.73 | 0.82 | 37 |
| accuracy | | | 0.88 | 100 |
| macro avg | 0.90 | 0.85 | 0.86 | 100 |
| weighted avg | 0.89 | 0.88 | 0.88 | 100 |

```
=== Performance Metrics ===
True Positives (TP): 27
True Negatives (TN): 61
False Positives (FP): 2
False Negatives (FN): 10
Accuracy: 0.8800
Error Rate: 0.1200
Precision: 0.9310
Recall: 0.7297
R² Score (using probabilities): 0.6407
```